

GPU Matrix multiplication

■ The `matmult_f.nvcc` driver is provided

```
matmult_f.nvcc type m n k [bs]
```

where `m`, `n`, `k` are the parameters defining the matrix sizes, `bs` is the optional blocksize for the block version, and `type` can be one of:

```
nat      - the native/naive version
lib      - the library version (note that this now calls a multithreaded
library)
gpu1     - the first gpu version
gpu2     - the second gpu version
gpu3     - the third gpu version
gpu4     - the fourth gpu version
gpu5     - the fifth gpu version
gpu6     - the sixth gpu version
gpulib   - the CUBLAS library version
```

as well as `blk`, `mnk`, `nmk`, ... (the permutations).

■ See README for more (also week 1 README)

GPU Matrix multiplication

- Driver uses `dlsym` for dynamic linking (C99 library) – therefore all functions in your shared library must have C naming convention
- Use `extern "C" {}`

```
extern "C" {  
    #include <cbblas.h> // The "C" headers also inside  
    ...  
    void matmult_lib(...)  
    {  
        ...  
    }  
    ...etc.  
}
```

GPU Matrix multiplication

- New cuBlas API used in the assignment
 - ❑ `#include <cublas_v2.h>`
- Creating a handle (to cuBlas context)
 - ❑ Important when using multiple host threads and multiple GPUs and makes it reentrant (non-blocking)
 - ❑ First creation has a large overhead!!

```
cublasStatus_t stat;  
cublasHandle_t handle;  
stat = cublasCreate(&handle);
```

- The `matmult_f.nvcc` driver creates a handle (and destroys it again) to "wake up" cuBlas

GPU Matrix multiplication

- For large matrices please use

- ❑ `MFLOPS_MAX_IT=1 ./matmult_f.nvcc ...`

- For benchmarks please use

- ❑ `MATMULT_COMPARE=0`

- For running on 1 core on CPU

- ❑ `MKL_NUM_THREADS=1`

- Overflow for large matrices (Please note!!)

- ❑ You should have the same overflow for all methods!

```
n-62-12-19(hhbs) $ ./matmult_f.nvcc gpulib 2048 2048 2048
98304.000 1178532.905 300 # matmult_gpulib
```




Overflow – not an actual error!

GPU Matrix multiplication


- Profiler (if available) may help to explain results

```
❑ matmult_f.nvcc gpu2 1024 1024 1024
```



Memory
bound!

```
❑ matmult_f.nvcc gpulib 1024 1024 1024
```

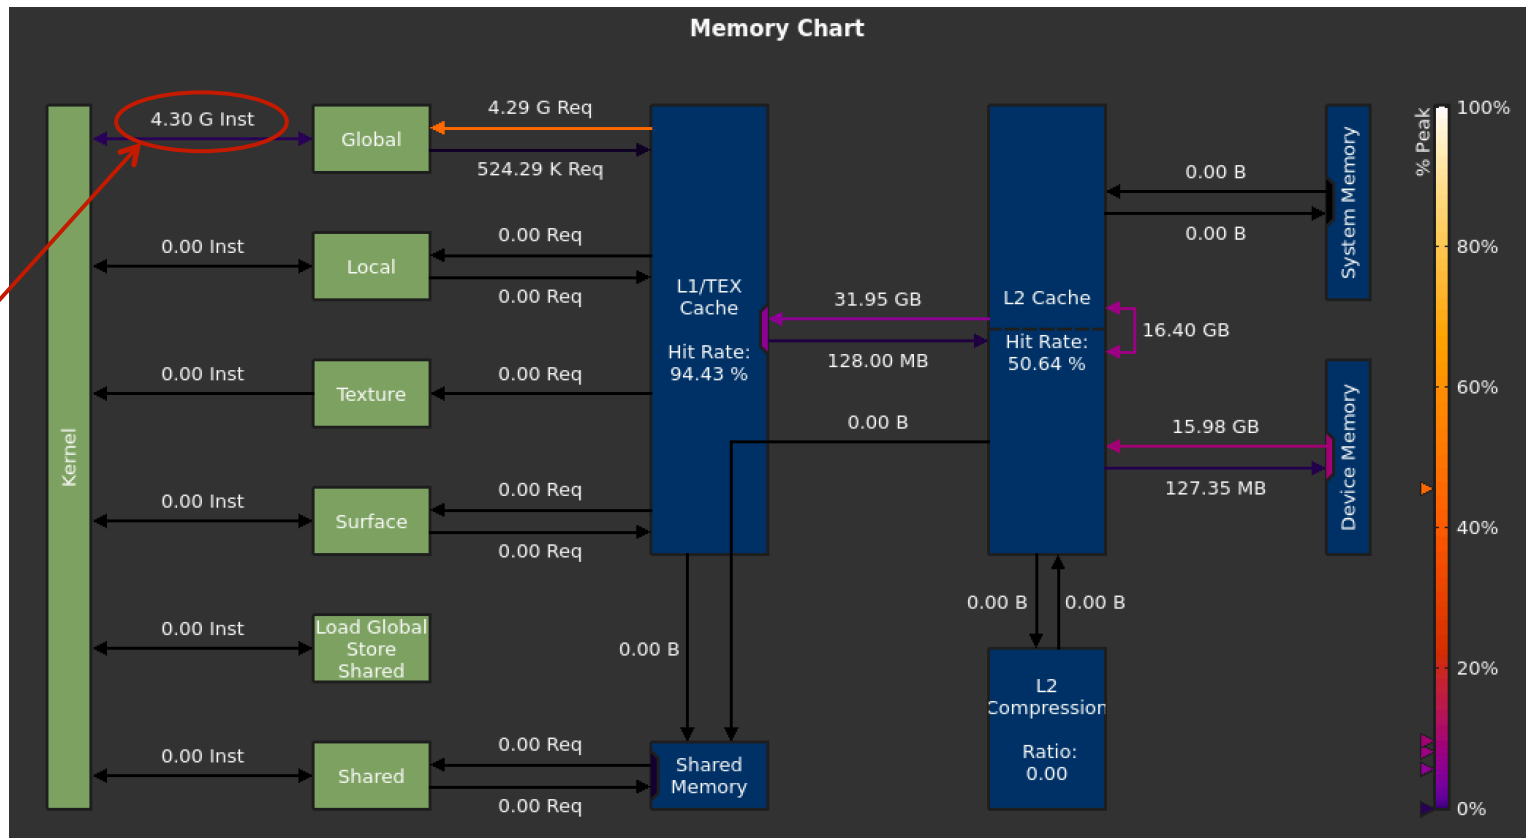


Compute
bound!

GPU Matrix multiplication

- Profiler (if available) may help to explain results


❑ `matmult_f.nvcc gpu2 4096 4096 4096`



GPU Poisson problem

- Multi-GPU version using Peer-to-peer access
 - Split task into two – top and bottom
 - Interior points can be updated from global memory
 - Border points must read "peer values" from other GPU

Read from global memory



Available from other GPU

Submitting GPU batch jobs

- Benchmark runs should always be submitted to queue `hpcintrogpu`
- See https://www.hpc.dtu.dk/?page_id=2759
- Maximum wall-clock time on jobs is 1 hour!
- For jobs using two GPUs use `num=2`
- For CPU-only jobs please do not request GPUs
- The `-G` flag sets debug lines into your code
 - ❑ This reduces the performance drastically
 - ❑ Remove it for performance tuning!!!

Reports / Analysis

- Assess your performance (Gflops / Bandwidth)
 - Profiler gives these numbers (“Throughput”)
- Speed-up calculations (fair)
 - If using 1 GPU, compare with 1 CPU!
 - `numactl --cpunodebind=0` to bind threads to cpu 0
- Tuning considerations
- Profiler analysis (if available)
- Relevant comments and observations
- Please keep the report format as close to the assignment questions as possible
- Please try to have one section in the report for each question in the assignment (helps us read it)